

**Name:** Mareena Fernandes

**Roll no.:** 8669

**Class:** TE IT

**Batch:** B

Experiment No.: 4

timestamp\_log.json:

```
{
  "log": [
    {
      "transaction": 25,
      "timestamp": 25,
      "operation": "read",
      "file": "B"
    },
    {
      "transaction": 26,
      "timestamp": 26,
      "operation": "read",
      "file": "B"
    },
    {
      "transaction": 26,
      "timestamp": 26,
      "operation": "write",
      "file": "B"
    },
    {
      "transaction": 25,
      "timestamp": 25,
      "operation": "read",
      "file": "A"
    },
    {
      "transaction": 26,
      "timestamp": 26,
      "operation": "read",
      "file": "A"
    }
  ]
}
```

```
{
  "transaction": 25,
  "timestamp": 25,
  "operation": "display",
  "file": "AB"
},
{
  "transaction": 26,
  "timestamp": 26,
  "operation": "write",
  "file": "A"
},
{
  "transaction": 26,
  "timestamp": 26,
  "operation": "display",
  "file": "AB"
}
]
```

exp4.py:

```
import json
from collections import defaultdict

file_map = {}
current_file = 0
files = []
transaction_map = {}
transactions = []

with open('./timestamp_log.json') as f:
    log = json.load(f)["log"]

    rts = [0] * len(log)
    wts = [0] * len(log)

    print("Timestamp Based Protocol Implementation: ")
    print("Op\t|\tFile\t|\tOutput")
    for data in log:
        if not data["file"] in files:
            file_map[data["file"]] = current_file
            current_file += 1
            files.append(data["file"])

        if not data["transaction"] in transactions:
            transaction_map[data["transaction"]] = data["timestamp"]
            transactions.append(data["transaction"])

        mapped_file_num = file_map[data["file"]]

        if data["operation"] == "read":
            if wts[mapped_file_num] > data["timestamp"]:
                transaction_num = list(transaction_map.keys())[list(
                    transaction_map.values()).index(wts[mapped_file_num])]

                print(
                    "Read\t|\t", data["file"], "\t|\tRollback and execute after Transaction", transaction_
num)

            elif data["timestamp"] > rts[mapped_file_num]:
                rts[mapped_file_num] = data["timestamp"]

                print("Read\t|\t", data["file"], "\t|\tExecute operation")
```

```

elif data["operation"] == "write":
    if rts[mapped_file_num] > data["timestamp"]:
        transaction_num = list(transaction_map.keys())[list(
            transaction_map.values()).index(wts[mapped_file_num])]

        print(
            "Write\t|\t", data["file"], "\t|\tRollback and execute after Transaction", transaction
_num)

    elif wts[mapped_file_num] > data["timestamp"]:
        print(
            "Write\t|\t", data["file"], "\t|\tReject and rollback (Obsolete write)")

    else:
        wts[mapped_file_num] = data["timestamp"]

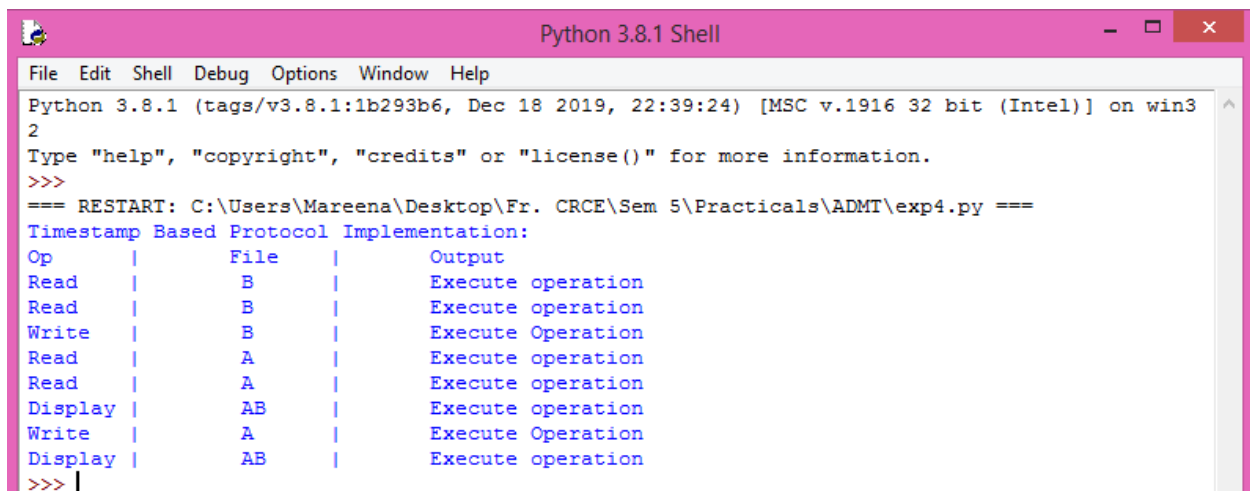
        print(
            "Write\t|\t", data["file"], "\t|\tExecute Operation")

elif data["operation"] == "display":
    print("Display\t|\t", data["file"], "\t|\tExecute operation")

else:
    print("Invalid operation type")

```

## Output:



```

Python 3.8.1 Shell
File Edit Shell Debug Options Window Help
Python 3.8.1 (tags/v3.8.1:1b293b6, Dec 18 2019, 22:39:24) [MSC v.1916 32 bit (Intel)] on win3
2
Type "help", "copyright", "credits" or "license()" for more information.
>>>
=== RESTART: C:\Users\Mareena\Desktop\Fr. CRCE\Sem 5\Practicals\ADMT\exp4.py ===
Timestamp Based Protocol Implementation:

```

Op	File	Output
Read	B	Execute operation
Read	B	Execute operation
Write	B	Execute Operation
Read	A	Execute operation
Read	A	Execute operation
Display	AB	Execute operation
Write	A	Execute Operation
Display	AB	Execute operation

```

>>>

```

## Post labs:

1. Consider the transactions T1, T2, and T3 and the schedules S1 and S2 given below

T1: r1(X); r1(Z); w1(X); w1(Z)

T2: r2(Y); r2(Z); w2(Z)

T3: r3(Y); r3(X); w3(Y)

S1: r1(X); r3(Y); r3(X); r2(Y); r2(Z);  
w3(Y); w2(Z); r1(Z); w1(X); w1(Z)

S2: r1(X); r3(Y); r2(Y); r3(X); r1(Z);  
r2(Z); w3(Y); w1(X); w2(Z); w1(Z)


Which one of the schedules S1, S2 is conflict serializable?

Ans:

Given Schedule S<sub>1</sub>:

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
r <sub>1</sub> (x)		
		r <sub>3</sub> (y)
		r <sub>3</sub> (x)
	r <sub>2</sub> (y)	
	r <sub>2</sub> (z)	
		w <sub>3</sub> (y)
	w <sub>2</sub> (z)	
r <sub>1</sub> (z)		
w <sub>1</sub> (x)		
w <sub>1</sub> (z)		

Precedence Graph:



Since no loop/ cycle exists in the precedence graph.

∴ Given schedule S<sub>1</sub> is conflict serializable.

To make Serial Schedule:


Indegree  $T_1 = 2$

Indegree  $T_2 = 0$

Indegree  $T_3 = 1$

$\therefore$  Remove  $T_2$  from graph and add it in sequence.

Precedence Graph:




Indegree  $T_1 = 1$

Indegree  $T_3 = 0$

$\therefore$  Remove  $T_3$  from graph and add it in sequence.

Precedence Graph:



Indegree  $T_1 = 0$

$\therefore$  Sequence for serial schedule is,

$[T_2 \rightarrow T_3 \rightarrow T_1]$


$T_1$	$T_2$	$T_3$
	$r_2(y)$	
	$r_2(z)$	
	$w_2(z)$	
		$r_3(y)$
		$r_3(x)$
		$w_3(y)$
$r_1(x)$		
$r_1(z)$		
$w_1(x)$		
$w_1(z)$		

Serial Schedule  $S_1'$

Given Schedule  $S_2$ :

$T_1$	$T_2$	$T_3$
$r_1(x)$		
		$r_3(y)$
	$r_2(y)$	
		$r_3(x)$
$r_1(z)$		
	$r_2(z)$	
		$w_3(y)$
$w_1(x)$		
	$w_2(z)$	
$w_1(z)$		

Precedence Graph:



Since there exists a cycle in the precedence graph.

$\therefore$  Given schedule  $S_2$  is not conflict serializable.

**2. List the disadvantages of time stamp protocol. And explain what is Thomas write rule.**

Ans:

- Recoverability:

In timestamp-based protocol, since the only limiting factor is the timestamp being the correct value, it is possible that a transaction reads data from an uncommitted write which a failure would be unrecoverable. This causes cascading rollback to that pointing time thus wasting time and computer power timestamp resolution. If minimum time elapsed between multiple transactions is less than resolution or coarseness of timestamp, it is possible to assign some timestamp to multiple transactions.

- Thomas Write Rule:

In timestamp-based protocol, if any operations don't follow order of serializability, it is rejected and rollback. Some operations can be harmless though and do not require rollback. Thomas Write Rule identifies updated writes to be harmless and thus instead of rolling back, simply ignores the operation moving on the next.

- a) If  $RTS(X) > TS(T_i)$ , then abort and rollback to ensure consistency.
- b) If  $WTS(X) > TS(T_i)$ , don't execute write operation and continue since even if performed before timestamp, it would be overwritten in future thus making it outdated or obsolete write.



## Experiment No.: 5

aries.json:

```
{
  "log": [
    {
      "lsn": 1,
      "operation": "write",
      "transaction": 1,
      "file": "A",
      "page": 1
    },
    {
      "lsn": 2,
      "operation": "write",
      "transaction": 2,
      "file": "B",
      "page": 1
    },
    {
      "lsn": 3,
      "operation": "write",
      "transaction": 2,
      "file": "C",
      "page": 2
    },
    {
      "lsn": 4,
      "operation": "flush",
      "page": 2
    },
    {
      "lsn": 5,
      "operation": "write",
      "transaction": 1,
      "file": "D",
      "page": 2
    },
    {
      "lsn": 6,
      "operation": "commit",
      "transaction": 1
    }
  ]
}
```

```
},  
{  
  "lsn": 7,  
  "operation": "write",  
  "transaction": 2,  
  "file": "B",  
  "page": 1  
},  
{  
  "lsn": 8,  
  "operation": "end",  
  "transaction": 1  
},  
{  
  "lsn": 9,  
  "operation": "crash"  
}  
]  
}
```

exp5.py:

```
import json
from collections import defaultdict

dirty_rec = defaultdict(lambda: 0)
trans_rec = defaultdict(lambda: 0)
page_rec = defaultdict(lambda: 0)

committed = []

undo_lsn = []
first_redo = float("inf")
last_undo = 0

with open('./aries.json') as f:
    log = json.load(f)["log"]

    print("Log: ")
    for data in log:
        print(data)
        if data["operation"] == "read" or data["operation"] == "write":
            trans_rec[data["transaction"]] = data["lsn"]
            if dirty_rec[data["page"]] == 0:
                dirty_rec[data["page"]] = data["lsn"]

        elif data["operation"] == "commit":
            if not data["transaction"] in committed:
                committed.append(data["transaction"])

        elif data["operation"] == "flush":
            dirty_rec[data["page"]] = 0
            page_rec[data["page"]] = data["lsn"]

        elif data["operation"] == "crash":
            break

        elif data["operation"] == "end":
            continue

    else:
        print("Invalid log file. Illegal operation type")
        break
```

```

print("\nTransaction Table: ")
print("TID\t|\tLSN\t|\tStatus")
for key, val in trans_rec.items():
    print(key, "\t|\t", val, end="\t|\t")
    if key in committed:
        print("Committed")
    else:
        if val > last_undo:
            last_undo = val
        print("Running")

print("\nDirty Page Table: ")
print("PageID\t|\tRecLSN")
for key, val in dirty_rec.items():
    if val != 0:
        print(key, "\t|\t", val)
        if val < first_redo:
            first_redo = val

print("\nRedo starts at LSN", first_redo)

for i in range(first_redo - 1, len(log)):
    data = log[i]
    if data["operation"] == "read" or data["operation"] == "write":
        if dirty_rec[data["page"]] > 0 and dirty_rec[data["page"]] <= data["lsn"] and page_rec[da
ta["page"]] < data["lsn"]:
            print("LSN", data["lsn"], ": Redo")
        else:
            print("LSN", data["lsn"], ": No Redo")
    elif data["operation"] == "crash":
        continue
    else:
        print("LSN", data["lsn"], ": Skip")

print("\nUndo starts up from LSN", last_undo,
      "for all Transaction IDs not in", committed)

for i in range(last_undo - 1, -1, -1):
    data = log[i]
    if data["operation"] == "read" or data["operation"] == "write":
        if not data["transaction"] in committed:
            undo_lsn.append(data["lsn"])

print("Undo LSN List: ", undo_lsn)

```

## Output:

```
Python 3.8.1 Shell
File Edit Shell Debug Options Window Help
Python 3.8.1 (tags/v3.8.1:1b293b6, Dec 18 2019, 22:39:24) [MSC v.1916 32 bit (Intel)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
=== RESTART: C:\Users\Mareena\Desktop\Fr. CRCE\Sem 5\Practicals\ADMT\exp5.py ===
Log:
{'lsn': 1, 'operation': 'write', 'transaction': 1, 'file': 'A', 'page': 1}
{'lsn': 2, 'operation': 'write', 'transaction': 2, 'file': 'B', 'page': 1}
{'lsn': 3, 'operation': 'write', 'transaction': 2, 'file': 'C', 'page': 2}
{'lsn': 4, 'operation': 'flush', 'page': 2}
{'lsn': 5, 'operation': 'write', 'transaction': 1, 'file': 'D', 'page': 2}
{'lsn': 6, 'operation': 'commit', 'transaction': 1}
{'lsn': 7, 'operation': 'write', 'transaction': 2, 'file': 'B', 'page': 1}
{'lsn': 8, 'operation': 'end', 'transaction': 1}
{'lsn': 9, 'operation': 'crash'}

Transaction Table:
TID      |      LSN      |      Status
1        |      5        |      Committed
2        |      7        |      Running

Dirty Page Table:
PageID   |      RecLSN
1        |      1
2        |      5

Redo starts at LSN 1
LSN 1 : Redo
LSN 2 : Redo
LSN 3 : No Redo
LSN 4 : Skip
LSN 5 : Redo
LSN 6 : Skip
LSN 7 : Redo
LSN 8 : Skip

Undo starts up from LSN 7 for all Transaction IDs not in [1]
Undo LSN List:  [7, 3, 2]
>>> |
```

## Post labs:

### 1. Implement ARIES algorithm for the above given sample log file.

Ans:

#### Analysis Phase:

LSN	T_id	Prev_LSN	Operation	Page_id
01	T <sub>1</sub>	-	Write(A)	P <sub>1</sub>
02	T <sub>2</sub>	-	Write(B)	P <sub>1</sub>
03	T <sub>3</sub>	02	Write(C)	P <sub>2</sub>
FLUSH P <sub>2</sub> TO DISK				
05	T <sub>1</sub>	01	Write(D)	P <sub>2</sub>
06	T <sub>1</sub>	05	Commit	-
07	T <sub>2</sub>	03	Write(B)	P <sub>1</sub>
08	T <sub>1</sub>	06	End	-
09	-	-	Crash	-

Page LSN = P<sub>2</sub> = 03

Transaction Table:

T_id	LSN	Status
T <sub>1</sub>	<del>1 5</del> 6	Committed
T <sub>2</sub>	<del>2 3</del> 7	Running

Dirty Page Table:

Page_id	Rec_LSN
P <sub>1</sub>	01
<del>P<sub>2</sub></del>	<del>03</del>
P <sub>2</sub>	05

← Removed

#### Redo Phase:

LSN 01 – Redo

{  
P<sub>1</sub> in DPT  
Rec LSN ≤ Current LSN  
Page LSN < Current LSN  
}

LSN 02 – Redo

$\left\{ \begin{array}{l} P_1 \text{ in DPT} \\ \text{Rec LSN} \leq \text{Current LSN} \\ \text{Page LSN} < \text{Current LSN} \end{array} \right\}$

LSN 03 – Redo

$\left\{ \begin{array}{l} P_2 \text{ in DPT} \\ \text{Rec LSN} \leq \text{Current LSN} \end{array} \right\}$

LSN 04 – Skip

{Flush}

LSN 05 – Redo

$\left\{ \begin{array}{l} P_2 \text{ in DPT} \\ \text{Rec LSN} \leq \text{Current LSN} \\ \text{Page LSN} < \text{Current LSN} \end{array} \right\}$

LSN 06 – Skip

{Commit}

LSN 07 – Redo

$\left\{ \begin{array}{l} P_1 \text{ in DPT} \\ \text{Rec LSN} \leq \text{Current LSN} \\ \text{Page LSN} < \text{Current LSN} \end{array} \right\}$

LSN 08 – Skip

{End}

LSN 09 – Skip

{Crash}

Undo Phase:

Transaction Table:

T_id	LSN	Status
T <sub>2</sub>	7	Running

Start Undo from LSN 7

LSN	T_id	Prev_LSN	Operation	Page_id
01	T <sub>1</sub>	-	Write(A)	P <sub>1</sub>
02	T <sub>2</sub>	-	Write(B)	P <sub>1</sub>
03	T <sub>2</sub>	02	Write(C)	P <sub>2</sub>
04	-	-	Flush P <sub>2</sub>	-
05	T <sub>1</sub>	01	Write(D)	P <sub>2</sub>
06	T <sub>1</sub>	05	Commit	-
07	T <sub>2</sub>	03	Write(B)	P <sub>1</sub>
08	T <sub>1</sub>	06	End	-
09	T <sub>2</sub>		Undo LSN: 07	P <sub>1</sub>
10	T <sub>2</sub>		Undo LSN: 03	P <sub>2</sub>
11	T <sub>2</sub>		Undo LSN: 02	P <sub>1</sub>



Experiment No.: 6

```
CREATE TABLE pay (Title VARCHAR(20) PRIMARY KEY, Salary NUMERIC(6));
```

```
INSERT INTO pay VALUES ('Elect. Engg.', 40000);
```

```
INSERT INTO pay VALUES ('Syst. Analy.', 50000);
```

```
INSERT INTO pay VALUES ('Mech. Engg.', 42000);
```

```
INSERT INTO pay VALUES ('Programmer', 65000);
```

```
SELECT * FROM pay;
```

```
CREATE TABLE emp (Eno NUMERIC(5) PRIMARY KEY, Ename VARCHAR(20), Title VARCHAR(20)  
REFERENCES pay);
```

```
INSERT INTO emp VALUES (101, 'John', 'Elect. Engg.');
```

```
INSERT INTO emp VALUES (102, 'Sam', 'Syst. Analy.');
```

```
INSERT INTO emp VALUES (103, 'Robert', 'Mech. Engg.');
```

```
INSERT INTO emp VALUES (104, 'Kim', 'Programmer');
```

```
INSERT INTO emp VALUES (105, 'Robert', 'Elect. Engg.');
```

```
INSERT INTO emp VALUES (106, 'Jack', 'Syst. Analy.');
```

```
INSERT INTO emp VALUES (107, 'Som', 'Programmer');
```

```
INSERT INTO emp VALUES (108, 'Smith', 'Elect. Engg.');
```

```
INSERT INTO emp VALUES (109, 'Albert', 'Mech. Engg.');
```

```
INSERT INTO emp VALUES (110, 'Bolt', 'Mech. Engg.');
```

```
SELECT * FROM emp;
```

### 1. Horizontal Fragmentation-

```
CREATE TABLE pay1 AS SELECT * FROM pay WHERE Salary<=45000;
```

```
CREATE TABLE pay2 AS SELECT * FROM pay WHERE Salary>45000;
```

```
SELECT * FROM pay1;
```

```
SELECT * FROM pay2;
```

```
SELECT * FROM pay1 UNION ALL SELECT * FROM pay2;
```

### 2. Vertical Fragmentation-

```
CREATE TABLE emp1 AS SELECT Eno, emp.Ename FROM emp;
```

```
SELECT * FROM emp1;
```

```
CREATE TABLE emp2 AS SELECT Eno, emp.Title FROM emp;
```

```
SELECT * FROM emp2;
```

```
SELECT emp1.Eno, emp1.Ename, emp2.Title FROM emp1, emp2 WHERE  
emp1.Eno=emp2.Eno;
```

```
SELECT emp1.Eno, emp1.Ename, emp2.Title FROM emp1 FULL OUTER JOIN emp2 ON  
emp1.Eno=emp2.Eno;
```

## Post labs:

### 1. Implement derived Hybrid Fragmentation.

Ans:

```
SELECT Eno, Ename FROM emp WHERE Eno>=105;
```

```
SELECT Eno, Title FROM emp WHERE Title='Elect. Engg.';
```

```
CREATE TABLE emp_frag1 AS SELECT emp.Eno, emp.Ename, emp.Title FROM emp, pay1 WHERE  
emp.Title=pay1.Title;
```

```
SELECT * FROM emp_frag1;
```

```
CREATE TABLE emp_frag2 AS SELECT emp.Eno, emp.Ename, emp.Title FROM emp, pay2 WHERE  
emp.Title=pay2.Title;
```

```
SELECT * FROM emp_frag2;
```

```
SELECT * FROM emp_frag1 UNION ALL SELECT * FROM emp_frag2;
```